

10601a: Homework #2 - “hello, Python/Java” (autolab)

TAs-in-charge:

Daniel Ribeiro Silva (drsilva@andrew)

Seth Flaxman (sflaxman@cs)

Assigned: Tuesday, 14 January 2014.

Due: 10am on Tuesday, 21 January 2014.

Late Penalty: 10% per day.

Policy on Collaboration among Students

The purpose of student collaboration is to facilitate learning, not to circumvent it. Studying the material in groups is strongly encouraged. It is also allowed to seek help from other students in understanding the material needed to solve a particular homework problem, provided no written notes are shared, or are taken at that time, and provided learning is facilitated, not circumvented. The actual solution must be done by each student alone, and the student should be ready to reproduce their solution upon request. The presence or absence of any form of help or collaboration, whether given or received, must be explicitly stated and disclosed in full by all involved, on the first page of their assignment. Specifically, **each assignment must contain a file named collaboration.txt where you will answer the following questions:**

- Did you receive any help whatsoever from anyone in solving this assignment? Yes / No. If you answered ‘yes’, give full details? (e.g. “Jane explained to me what is asked in Question 3.4”).
- Did you give any help whatsoever to anyone in solving this assignment? Yes / No. If you answered ‘yes’, give full details? (e.g. “I pointed Joe to section 2.3 to help him with Question 2”).

Collaboration without full disclosure will be handled severely, in compliance with CMU’s Policy on Cheating and Plagiarism. All violations (even first one) will carry severe penalties, up to failure in the course, and will in addition always be reported to the university authorities.

Some of the homework assignments used in this class may have been used in prior versions of this class, or in classes at other institutions. Avoiding the use of heavily tested assignments will detract from the main purpose of these assignments, which is to reinforce the material and stimulate think-

ing. Because some of these assignments may have been used before, solutions to them may be (or may have been) available online, or from other people. It is explicitly forbidden to use any such sources, or to consult people who have solved these problems before. You must solve the homework assignments completely on your own. I will strictly enforce this policy, and if a violation is detected it will be dealt with harshly. Collaboration with other students who are currently taking the class is allowed, but only under the conditions stated above.

General Instructions

The goal of this assignment is to introduce you to Java or Python as powerful tools for processing text and to introduce you to the autolab environment for homework submission. We've described functions that we want you to implement. Be sure to test your code before submitting.

In this homework you have to choose Python or Java as your programming language, and you must respect that choice for all questions in this homework. The autolab environment has Python 2.6 installed—check your code runs on `linux.andrew.cmu.edu`.

0 INTRODUCTION

Bag of Words (BoW) is a widely used method in machine learning and natural language processing for representing text information. The idea is to represent a given document by the count of the frequency of each word in the document—a so-called “bag of words.”

We illustrate with an example. Consider the following three documents:

Document1

come on you painter, you piper, you prisoner, and shine

Document2

The piper at the gates of dawn

Document3

Then the piper will lead us to reason.

There are 19 unique words that appear in the documents:

come, on, you, painter, piper, prisoner, and, shine, at, the, gates, of, dawn,
then, will, lead, us, to, reason

These 19 words are our vocabulary. Our BoW representation will be a 19-dimensional vector, one for each document. Every dimension of this vector represents the term frequency of a word in the vocabulary.

Thus, the BoW representation of the three documents will be:

	come	on	you	painter	piiper	prisoner	and	shine	at	the	gates	of	dawn	then	will	lead	us	to	reason
document1	1	1	3	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0
document2	0	0	0	0	1	0	0	0	1	2	1	1	1	0	0	0	0	0	0
document3	0	0	0	0	1	0	0	0	0	1	0	0	0	1	1	1	1	1	1

Now, instead of raw text, each document is represented by a feature vector. In this assignment you will build a simplified version of a BoW.

1 SPLITTING TEXT

Write a program that takes the location path of a text file as its argument and generates an output with a list of unique, lower case words, sorted in alphabetical order. You should only split the input text on spaces (' '). Don't worry about tabs, line breaks or any kind of punctuation marks. The input will be a very simple sentence.

For example, if the input text file contains the following sentence

The quick brown fox jumps over the lazy dog

your output should be

brown,dog,fox,jumps,lazy,over,quick,the

Your output is **not** supposed to contain a line break ('\n') at the end. Please observe that there are no spaces between words on the output.

You have to output your answer to the standard output (stdout).

We will run your code as below:

For Python:

```
python question1.py "path/to/input"
```

For Java:

```
javac question1.java
java question1 "path/to/input"
```

Make sure your program is called question1.py or question1.java.

2 COUNTING WORDS

Write a program that takes the location path of a text file as argument, splits the text as in question 1, and counts the number of occurrences of each word. The words must also be displayed in alphabetical order.

For example, if the input text file contains the following text:

```
The quick brown fox jumps over the lazy dog
```

your output should be

```
brown:1,dog:1,fox:1,jumps:1,lazy:1,over:1,quick:1,the:2
```

Your output is **not** supposed to contain a line break ('\n') at the end. Please observe that there are no spaces between words or numbers on the output.

You have to output your answer to the standard output (stdout).

We will run your code as below:

For Python:

```
python question2.py "path/to/input"
```

For Java:

```
javac question2.java  
java question2 "path/to/input"
```

Make sure your program is called question2.py or question2.java.

3 REMOVING STOP WORDS

Alter your program from question 2 so that it now removes stop words from a list provided. Stop words are words such as “the” or “at” that carry little meaning—they are usually removed during text processing for learning tasks.

For example, if the input text file contains the following text

```
The quick brown fox jumps over the lazy dog
```

and the stop words list is

as
the
in
over
on

your output should be

```
brown:1,dog:1,fox:1,jumps:1,lazy:1,quick:1
```

Your output file is **not** supposed to contain a line break ('\n') at the end. Please observe that there are no spaces between words or numbers on the output.

You have to output your answer to the standard output (stdout).

We will run your code as below:

For Python:

```
python question3.py "path/to/input/file" "path/to/stopwords/file"
```

For Java:

```
javac question3.java  
java question3 "path/to/input/file" "path/to/stopwords/file"
```

Make sure your program is called question3.py or question3.java.

4 AUTOLAB SUBMISSION

For Java:

You must submit a .tgz file named **hw2.tgz** containing all your files (the three .java files + collaboration.txt).

You can create that file by running “ *tar -cvf hw2.tgz *.java collaboration.txt*”.

DO NOT put the above files in a folder and then tar gzip the folder.

You must submit this file to the “**homework2java**” link on Autolab (<https://autolab.cs.cmu.edu/10601a-s14/homework2java/>). Since you have a limited number of attempts, please make sure your code is working before you submit it to Autolab.

For Python:

You must submit a .tgz file named **hw2.tgz** containing all your files (the three .py files + collaboration.txt).

You can create that file by running “ `tar -cvf hw2.tgz *.py collaboration.txt`”.

DO NOT put the above files in a folder and then tar gzip the folder.

You must submit this file to the “**homework2python**” link on Autolab (<https://autolab.cs.cmu.edu/10601a-s14/homework2python/>). Since you have a limited number of attempts, please make sure your code is working before you submit it to Autolab.